

Introduction to MicroBlaze

LECTURE 5

IOULIIA SKLIAROVA

MicroBlaze Processor

Soft processor

- ~1 900-7 000 logic cells (~1 200-4 500 LUTs) - estimates
- 63 400 LUTs available in Artix-7 XC7A100T

RISC architecture

- utilizes a small, highly-optimized set of instructions, rather than a more specialized set of instructions often found in other types of architectures

32/64-bit architecture

- Thirty-two 32-bit or 64-bit general purpose registers

In production since 2002

Supported in

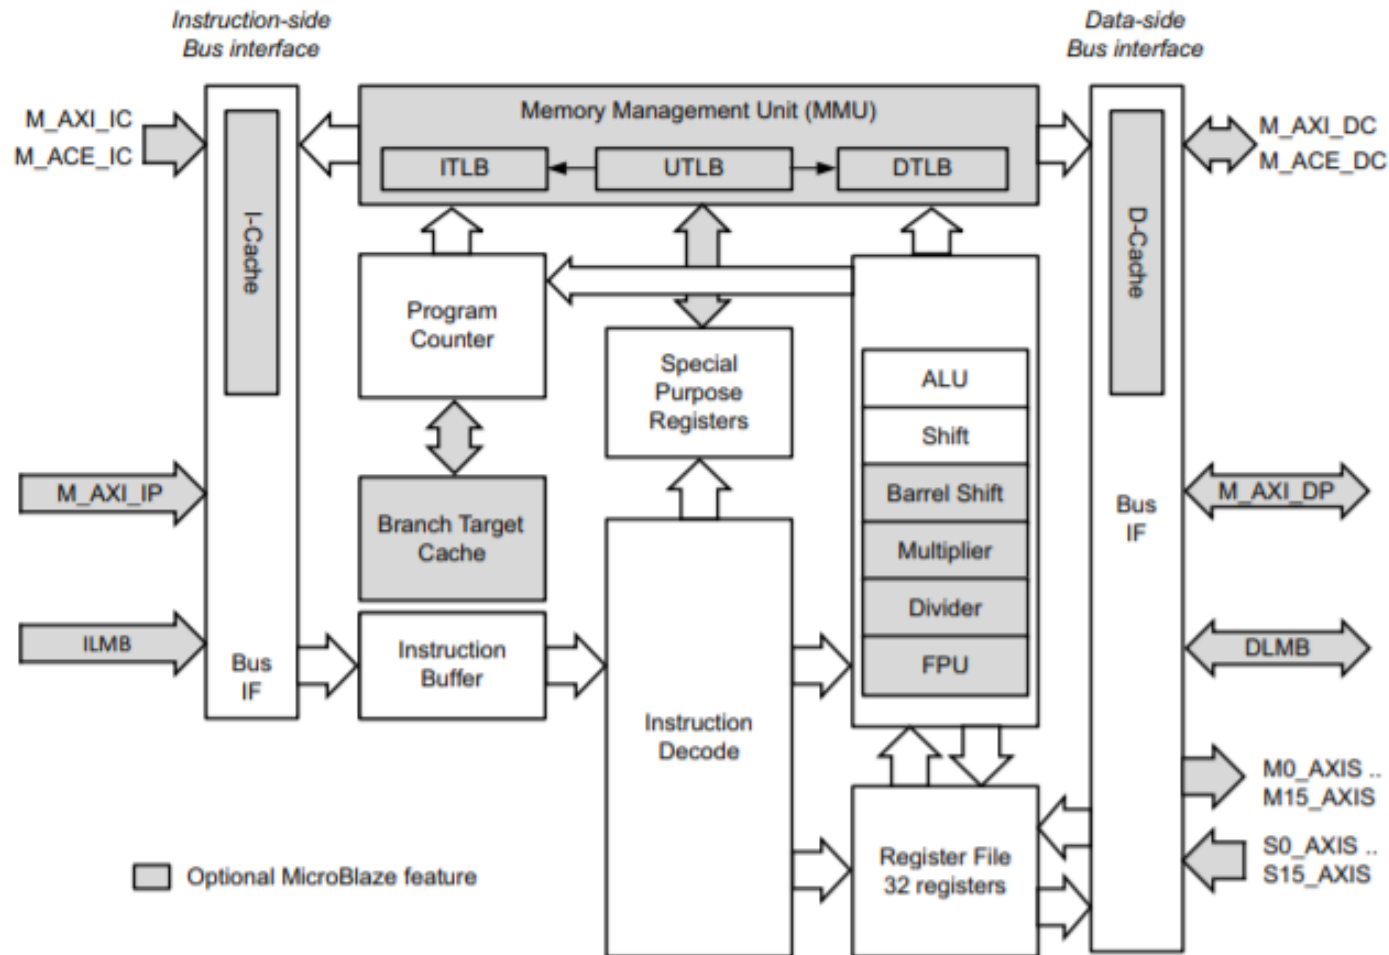
- 7-series/UltraScale/UltraScale+ devices

Three preset configurations:

- a simple microcontroller running bare-metal applications (~200 MHz in Artix-7);
- a real-time processor running FreeRTOS (~170 MHz);
- an application processor with a memory management unit running Linux (~140 MHz)

MicroBlaze Overview

MicroBlaze Processor Reference Guide - UG984



MicroBlaze Features

The fixed feature set of the processor includes:

- Thirty-two 32-bit or 64-bit general purpose registers
- 32-bit instruction word with three operands and two addressing modes
- Default 32-bit address bus, extensible to 64 bits
- Single issue pipeline

Configurable features

- Processor pipeline depth
- Floating-point unit (FPU)
- Hardware divider
- Area or speed optimized
- 64-bit mode
- ...

MicroBlaze Data Types

The MicroBlaze processor uses Big-Endian or Little-Endian (default) format to represent data, depending on the selected endianness.

The hardware supported data types for 32-bit MicroBlaze are word, half word, and byte. With 64-bit MicroBlaze the data types long and double are also available in hardware.

| | | | | |
|-----------------------------------|--------|-----|-----|--------|
| Big-Endian Byte Address | n | n+1 | n+2 | n+3 |
| Big-Endian Byte Significance | MSByte | | | LSByte |
| Big-Endian Byte Order | n | n+1 | n+2 | n+3 |
| Big-Endian Byte-Reversed Order | n+3 | n+2 | n+1 | n |
| Little-Endian Byte Address | n+3 | n+2 | n+1 | n |
| Little-Endian Byte Significance | MSByte | | | LSByte |
| Little-Endian Byte Order | n+3 | n+2 | n+1 | n |
| Little-Endian Byte-Reversed Order | n | n+1 | n+2 | n+3 |
| Bit Label | 0 | | | 31 |
| Bit Significance | MSBit | | | LSBit |

MicroBlaze Instruction Summary

All MicroBlaze instructions are 32 bits and are defined as either Type A or Type B.

Type A instructions have up to two source register operands and one destination register operand.

Type B instructions have one source register and a 16-bit immediate operand. Type B instructions have a single destination register operand.

Table 2-7: MicroBlaze Instruction Set Summary

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|----------------|--------|------|-------|-------|-------------|--|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| ADD Rd,Ra,Rb | 000000 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + Ra$ |
| RSUB Rd,Ra,Rb | 000001 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + \overline{Ra} + 1$ |
| ADDC Rd,Ra,Rb | 000010 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + Ra + C$ |
| RSUBC Rd,Ra,Rb | 000011 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + \overline{Ra} + C$ |
| ADDK Rd,Ra,Rb | 000100 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + Ra$ |
| RSUBK Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 00L00000000 | $Rd := Rb + \overline{Ra} + 1$ |
| CMP Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 00L00000001 | $Rd := Rb + \overline{Ra} + 1$ $Rd[0] := 0$ if $(Rb \geq Ra)$ else $Rd[0] := 1$ |
| CMPU Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 00L00000011 | $Rd := Rb + \overline{Ra} + 1$ (unsigned) $Rd[0] := 0$ if $(Rb \geq Ra, \text{unsigned})$ else $Rd[0] := 1$ |
| | | | | | | |

MicroBlaze Pipeline

MicroBlaze instruction execution is pipelined.

For most instructions, each stage takes one clock cycle to complete.

Consequently, the number of clock cycles necessary for a specific instruction to complete is equal to the number of pipeline stages, and one instruction is completed on every cycle in the absence of data, control or structural hazards.

- A data hazard occurs when the result of an instruction is needed by a subsequent instruction. This can result in stalling the pipeline, unless the result can be forwarded to the subsequent instruction. The MicroBlaze GNU Compiler attempts to avoid data hazards by reordering instructions during optimization.
- A control hazard occurs when a branch is taken, and the next instruction is not immediately available. This results in stalling the pipeline. MicroBlaze provides delay slot branches and the optional branch target cache to reduce the number of stall cycles.
- A structural hazard occurs for a few instructions that require multiple clock cycles in the execute stage or a later stage to complete. This is achieved by stalling the pipeline.

Three Stage Pipeline

With the MicroBlaze is optimized for **area**, the pipeline is divided into three stages to minimize hardware cost: Fetch, Decode, and Execute.

The three stage pipeline does not have any data hazards.

Pipeline stalls are caused by control hazards, structural hazards due to multi-cycle instructions, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

| | cycle1 | cycle2 | cycle3 | cycle4 | cycle5 | cycle6 | cycle7 |
|---------------|--------|--------|---------|---------|---------|---------|---------|
| instruction 1 | Fetch | Decode | Execute | | | | |
| instruction 2 | | Fetch | Decode | Execute | Execute | Execute | |
| instruction 3 | | | Fetch | Decode | Stall | Stall | Execute |

Five Stage Pipeline

With the MicroBlaze is optimized for **performance**, the pipeline is divided into five stages to maximize performance: Fetch (IF), Decode (OF), Execute (EX), Access Memory (MEM), and Writeback (WB).

Pipeline stalls are caused by data hazards, control hazards, structural hazards due to multicycle instructions, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

| | cycle1 | cycle2 | cycle3 | cycle4 | cycle5 | cycle6 | cycle7 | cycle8 | cycle9 |
|---------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| instruction 1 | IF | OF | EX | MEM | WB | | | | |
| instruction 2 | | IF | OF | EX | MEM | MEM | MEM | WB | |
| instruction 3 | | | IF | OF | EX | Stall | Stall | MEM | WB |

Eight Stage Pipeline

With the MicroBlaze is optimized for **frequency**, the pipeline is divided into eight stages to maximize possible frequency: Fetch (IF), Decode (OF), Execute (EX), Access Memory 0 (M0), Access Memory 1 (M1), Access Memory 2 (M2), Access Memory 3 (M3) and Writeback (WB).

Pipeline stalls are caused by data hazards, control hazards, structural hazards, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

| | cycle1 | cycle2 | cycle3 | cycle4 | cycle5 | cycle6 | cycle7 | cycle8 | cycle9 | cycle10 | cycle11 |
|---------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|---------|---------|
| instruction 1 | IF | OF | EX | M0 | M1 | M2 | M3 | WB | | | |
| instruction 2 | | IF | OF | EX | M0 | M0 | M1 | M2 | M3 | WB | |
| instruction 3 | | | IF | OF | EX | Stall | M0 | M1 | M2 | M3 | WB |

Memory Architecture

MicroBlaze is implemented with a **Harvard memory architecture**; instruction and data accesses are done in separate address spaces.

- The **instruction address space** has a 32-bit virtual address range with 32-bit MicroBlaze (that is, handles up to 4GB of instructions), and can be extended up to a 64-bit physical address range.
- The **data address space** has a default 32-bit range, and can be extended up to a 64-bit range (that is, handles from 4GB to 16EB of data).

The instruction and data memory ranges can be made to overlap by mapping them both to the same physical memory. The latter is necessary for software debugging.

Both instruction and data interfaces of MicroBlaze are default 32 bits wide and use big endian or little endian, bit-reversed format, depending on the selected endianness. MicroBlaze supports word, halfword, and byte accesses to data memory.

Data accesses must be aligned (word accesses must be on word boundaries, halfword on halfword boundaries), unless the processor is configured to support unaligned exceptions. All instruction accesses must be word aligned.

Final Remarks

At the end of this lecture you should be able to:

- have a generic idea of MicroBlaze processor

To do:

- Complete the lab. 4