

Outline

- 1 Introduction
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Evolutionary Learning
- 5 Reinforcement Learning



Motivation

Programming robots is a hard work!

- No high-level programming language;
- Sensors and actuators are noisy;
- Robotics is moving towards increasingly unstructured environments.

If only robots could learn how to perform tasks by themselves. . .

Machine Learning

- Machine Learning is:
 - “A **computer program** is said to **learn** from **experience** E with respect to some **class of tasks** T and **performance measure** P, if its performance at tasks in T, as measured by P, **improves with experience** E” *Tom Mitchell. Machine Learning, 1997*
- Key concepts:
 - Experience (data);
 - Task;
 - Performance Measure (metric);
 - Improvement
- Machine Learning can be seen as a search problem of finding a policy that maps states to responses, $(\pi : S \rightarrow R)$, to perform a desired task.

Machine Learning




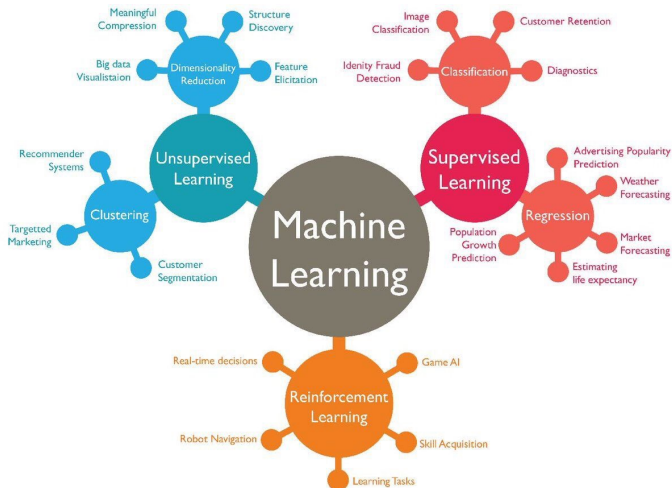
Figure: Data Science Venn Diagram, by Drew Conway

Machine Learning Paradigms

- Supervised Learning
 - Learn from examples
- Unsupervised Learning
 - Discover structure in data
- Reinforcement learning
 - Learn from interaction



Machine Learning Paradigms



Outline

- 1 Introduction
- 2 Supervised Learning**
- 3 Unsupervised Learning
- 4 Evolutionary Learning
- 5 Reinforcement Learning



Supervised Learning

- A “teacher” **provides** training data consisting of **states** (S) and the **desired response** (R).
 - The learning process knows the desired output for several inputs.
 - The supervisor indicates what is the best action to take (sometimes).
- The robot must learn to **fit** the training data data and **generalize** a policy to the states not covered by the training data.
- Common methods:
 - k Nearest Neighbor, Logistic Regression, Neural Networks, Decision Trees, Support Vector Machines, Gaussian Processes, ...

Supervised Learning

Label: 1



Label: 0




Label: 1




Label: 4




0




7



3



5



Supervised Learning




Figure: A regression problem: bias vs. variance

Supervised Learning

IRIS dataset

- Find flower species from:

- Sepal width
- Sepal length
- Petal width
- Petal length


iris setosa



petal

sepal

iris versicolor



petal

sepal

iris virginica




petal

sepal

Supervised Learning

IRIS dataset


- Considering only:
 - Sepal width
 - Sepal length



Supervised Learning

IRIS dataset - SVM


- Support Vector Machine Classification
 - Find the maximum margin hyperplane that separates categories



Supervised Learning

IRIS dataset - KNN


- K-Nearest Neighbor Classification
 - Find the majority among closest neighbors



Supervised Learning

IRIS dataset - Decision Tree

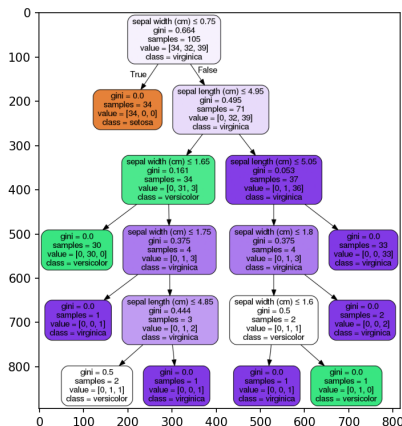
- Decision Tree Classification
 - Find best splitting rules based on feature values



Supervised Learning

IRIS dataset - Decision Tree

- Decision Tree Classification
 - Find best splitting rules based on feature values



Applications

- Very powerful when applied for Perception in Robotics!
 - Pattern Recognition for Robot Vision.
 - Probabilistic Models for Kalman and Particle Filters.
 - Fault detection.
 - Change detection.
- Learning from Demonstration for Control

Supervised Learning

Examples




Figure: ImageNet Visual Recognition Challenge Results.

Outline

- 1 Introduction
- 2 Supervised Learning
- 3 Unsupervised Learning**
- 4 Evolutionary Learning
- 5 Reinforcement Learning



Clustering




Figure: Clustering flowers.

Outline

- 1 Introduction
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Evolutionary Learning**
- 5 Reinforcement Learning



Evolutionary Learning

- Unsupervised/Semi-supervised Learning Paradigm.
- A **policy may be encoded** in **strings** (Genetic Algorithms) or in **computer programs** (Genetic Programming).
- The learning process works on a **set of policies** (generation).
- The robot is provided a **fitness function**.
- Each individual on the generation is evaluated given the fitness function.
- **Genetic operators**: selection, crossover, mutation.
- Generations are **regenerated** trying to get better fitness values

Evolutionary Learning

Examples

- Robot Motion
 - Applied in-house to learn biped walking gaits (Picado-2009).
 - Applied in-house to learn kick behavior (Abdolmaleki-2016).
 - Applied to learn the model of the robot (Lipson-2008).
- Robot Hardware
 - Applied to evolve the shape of the robot, to obtain previously unknown robot shapes (Lipson-2006).



Figure: One of Hod Lipson morphologically evolved robots.

Outline

- 1 Introduction
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Evolutionary Learning
- 5 Reinforcement Learning**



Reinforcement Learning




Figure: A Reinforcement Learning System.

- Unsupervised/Semi-supervised Learning Paradigm.
- Modeled as a Markov Decision Process:
 - 1 a set of states S ;
 - 2 a set of actions $A(s)$;
 - 3 a state transition model: $P(s, a, s') \rightarrow [0, 1]$
 - 4 a reward function : $R(s, a, s') \rightarrow r_t \in \mathbb{R}$
- The goal?
- To determine a policy that maximizes the return (total reward),
$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$
- How?
- By calculating a *Value Function*

Reinforcement Learning

- OpenAI gym (Markov Decision Process API)

```
import gym

# load the environment
env = gym.make('Example')

# perform N episodes
while True:

    # prepare the environment for the next episode
    obs = env.reset()

    # flag for episode completion
    done = False

    # run the episode
    while not done:

        # choose how to act based on the current state
        # usually the output of a neural network
        action = env.action_space.sample()

        # advance the simulation by one timestep
        # by interacting with the world
        obs, reward, done, info = env.step(action)

# cleanup
env.close()
```



Reinforcement Learning




Figure: A maze environment.

Reinforcement Learning

- A Value Function estimates the expected return for *all* states
- $V^{\pi}(s_t) = \mathbb{E}[R_t]$




Figure: The optimal value function.

Reinforcement Learning

- A *Value Function* represents “how good” it is to be in a given state, $V^\pi(s_t) : s_t \rightarrow \mathbb{E}[R_t]$
- Bellman Equation:
$$V^\pi(s) = r(s, \pi(s), s') + \gamma V(s'), s' = f(s, \pi(s))$$
- How can we learn a Value Function?

Reinforcement Learning

Algorithm 1 Policy Iteration (Policy Evaluation)

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$

```
1: repeat
2:   repeat
3:      $\Delta \leftarrow 0$ 
4:     for all  $s \in S$  do
5:        $v \leftarrow V(s)$ 
6:        $V(s) \leftarrow \sum_{s'} P(s, \pi(s), s') [r(s, \pi(s), s') + \gamma V(s')]$ 
7:        $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
8:     end for
9:   until  $\Delta \leq \theta$  ▷ a small positive value
10:  ...
11: until StablePolicy == True
```

Reinforcement Learning

Algorithm 2 Policy Iteration (Policy Improvement)

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$

```
1: repeat
2:   ...
3:   StablePolicy  $\leftarrow$  True
4:   for all  $s \in S$  do
5:      $b \leftarrow \pi(s)$ 
6:      $\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s, a, s') [r(s, a, s') + \gamma V(s')]$ 
7:     if  $b \neq \pi(s)$  then
8:       StablePolicy  $\leftarrow$  False
9:     end if
10:  end for
11: until StablePolicy == True
```

Reinforcement Learning

Algorithm 3 Value Iteration

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$

```
1: repeat
2:    $\Delta \leftarrow 0$ 
3:   for all  $s \in S$  do
4:      $v \leftarrow V(s)$ 
5:      $V(s) \leftarrow \max_a \sum_{s'} P(s, a, s') [r(s, a, s') + \gamma V(s')]$ 
6:      $\Delta \leftarrow \max(\Delta, \|v - V(s)\|)$ 
7:   end for
8: until  $\Delta \leq \theta$ 
```

▷ a small positive value

Reinforcement Learning

- That's nice, but what do we do with a *Value Function*?
- Extract a *policy*!
- A *policy* maps states to actions, $\pi : s \rightarrow a$
- $\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') [r(s, a, s') + \gamma V(s')]$

Reinforcement Learning




Figure: The optimal policy

Reinforcement Learning

- 1 Finite State Space
- 2 Finite Action Space
- 3 Curse of Dimensionality



Reinforcement Learning

- 1 Finite State Space
- 2 Finite Action Space
- 3 Curse of Dimensionality

Question?

What if we don't know the state transition model?

Example

- Where is the opponent dribbling the ball?
- To where in the goal is the opponent shooting the ball?

Reinforcement Learning

- 1 Finite State Space
- 2 Finite Action Space
- 3 Curse of Dimensionality

Question?

What if we don't know the state transition model?

Example

- Where is the opponent dribbling the ball?
- To where in the goal is the opponent shooting the ball?

Reinforcement Learning

Algorithm 4 The Q-Learning algorithm

Require: $Q(s, a)$ initialized arbitrarily $\forall s \in S, \forall a \in A(s)$

1: **loop**

2: Initialize $s = s_0$

3: **repeat**

4: $a \leftarrow \pi(s)$

5: take action a ; observe reward, r , and successor state s'

6: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_b Q(s', b) - Q(s, a)]$

7: $s \leftarrow s'$

8: **until** s is terminal

9: **end loop**

Reinforcement Learning




Figure: The optimal Q-function.

- $\pi(s) = \arg \max_a Q(s, a)$

Reinforcement Learning

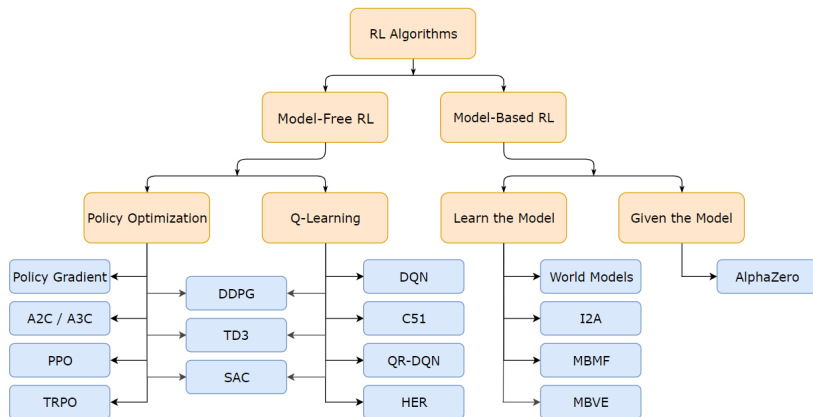


Figure: Reinforcement Learning Taxonomy.

Summary

- Machine Learning is a valid software development tool for programming robotic agents.
- Many paradigms, many challenges, many solutions, even more problems...
- The future of Robotics?



References

- Pattern Recognition and Machine Learning, Christopher Bishop, Springer, 2006
- Machine Learning, Tom Mitchell, McGraw Hill, 1997
- Reinforcement Learning: An Introduction, Richard Sutton, Andrew Barto, MIT Press, 2018
- OpenAI, scikit-learn, TensorFlow, PyTorch, stable baselines, Keras, RapidMiner, clsquare, OpenCV, Shark, ...